Step 1: After logging in with your credentials at http://192.168.12.1:31380, go to “Notebook Servers” on the left sidebar

Step 2: Select “New Server”

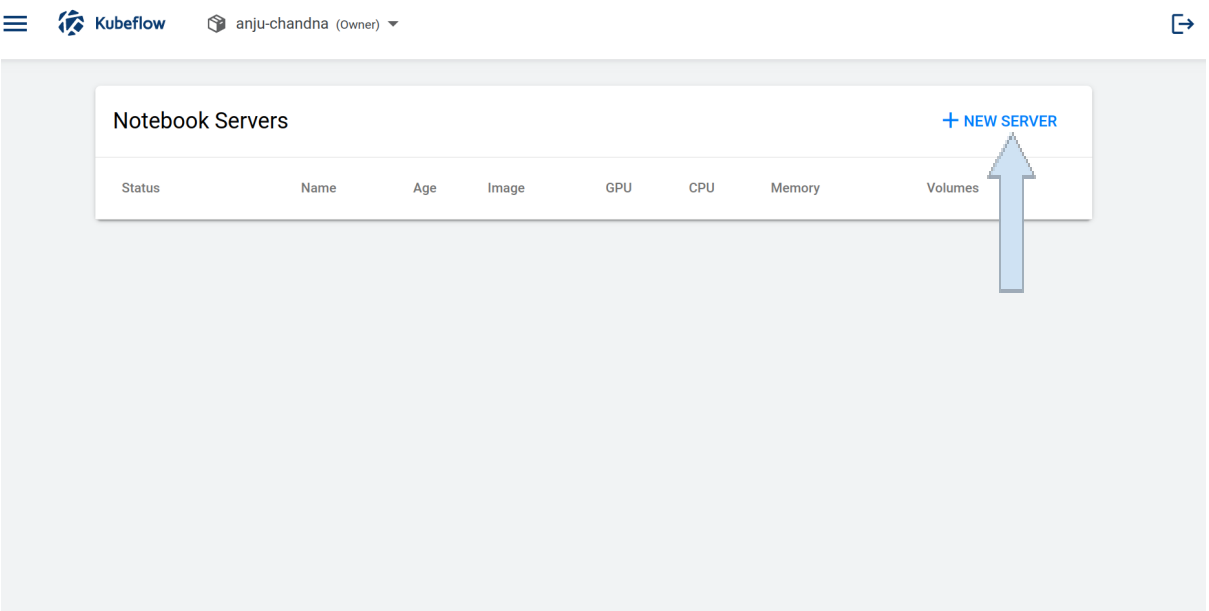

Step 3: Give a name to the notebook server and select an image of the ML framework

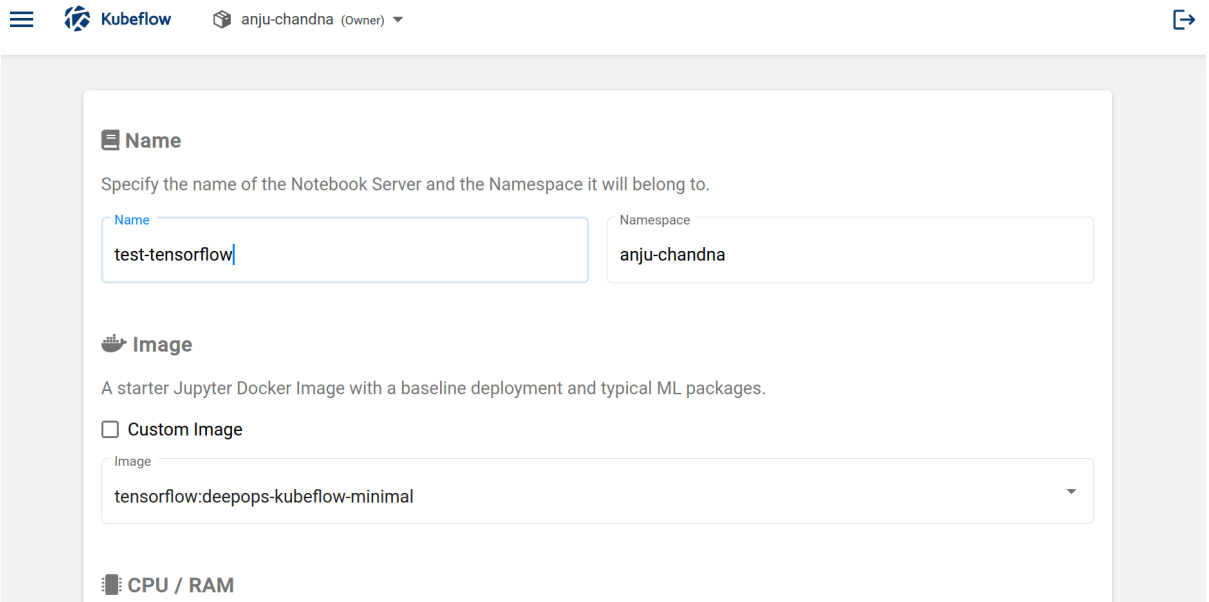

Step 4: Specify CPU, RAM, and GPU requirements (do not exceed your allotted quota)

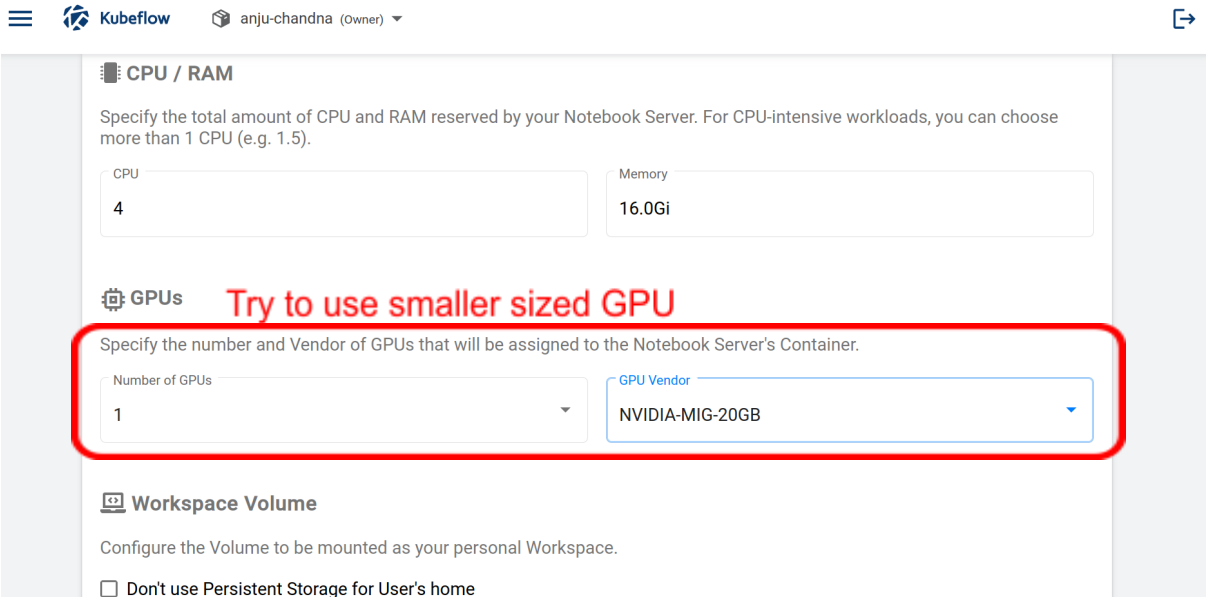

Step 5: Choose whether to use persistent storage or not

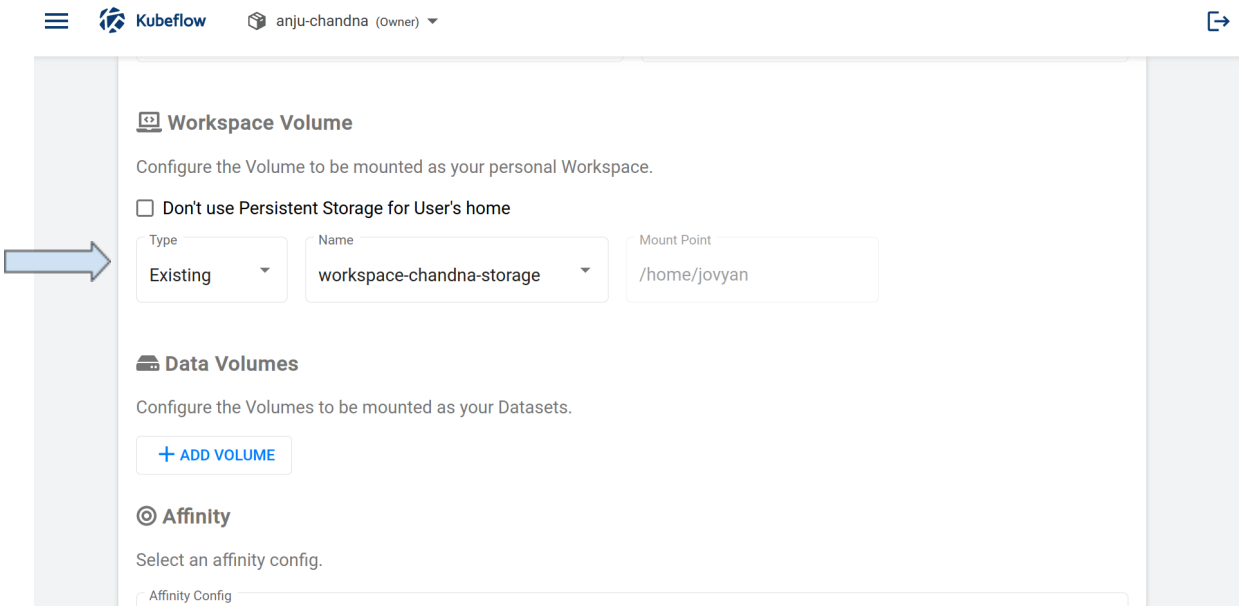

Do note that, you can create **up to three persistent storage volumes** and all the contents that you will store in **"/home/jovyan"** directory will be permanent and will persist even after you shut down the notebook server.

After you have created three persistent storage volumes, you have to reuse them for future notebook servers. When you want to re-use persistent storage please select “Type” as “Existing” and “Name” as your volume name (e.g., **"workspace-user-storage"**).

You can also choose to not use persistent storage and your contents will be deleted when you delete the notebook server.

You do not have to change the “Data Volumes” and “Affinity” settings.

Step 6: You do not have to change the “Tolerations” and “Configurations” settings and you can keep the “Enable Shared Memory” option selected.

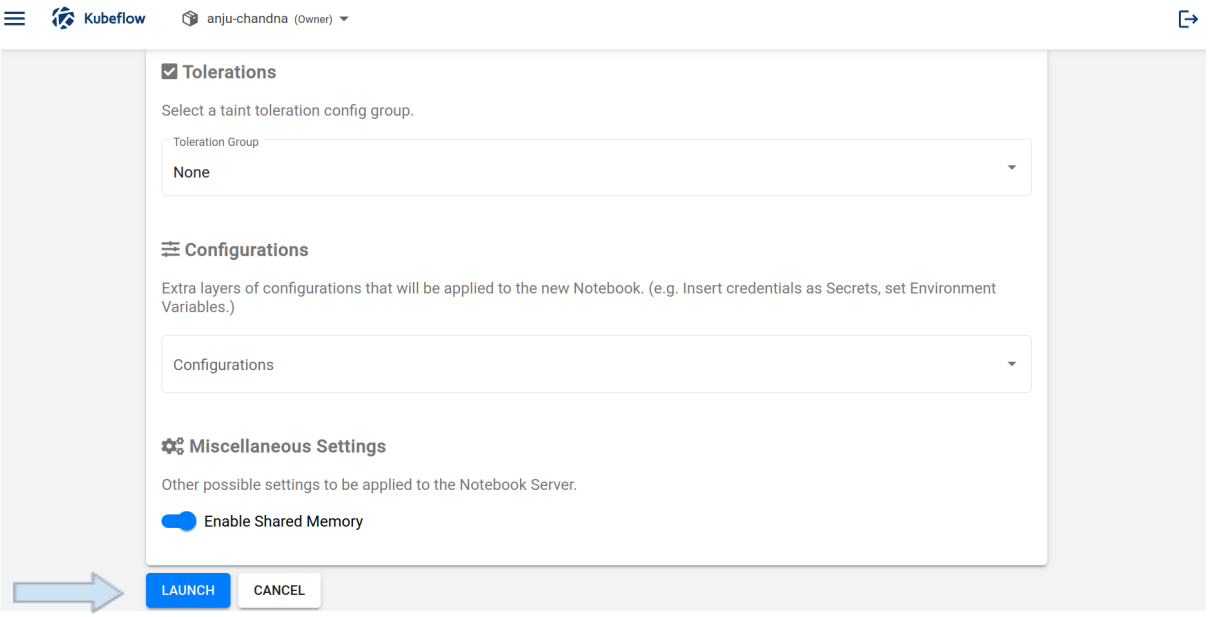

Step 7: If notebook creation is successful then you will see a green tick in status and then you connect to the notebook

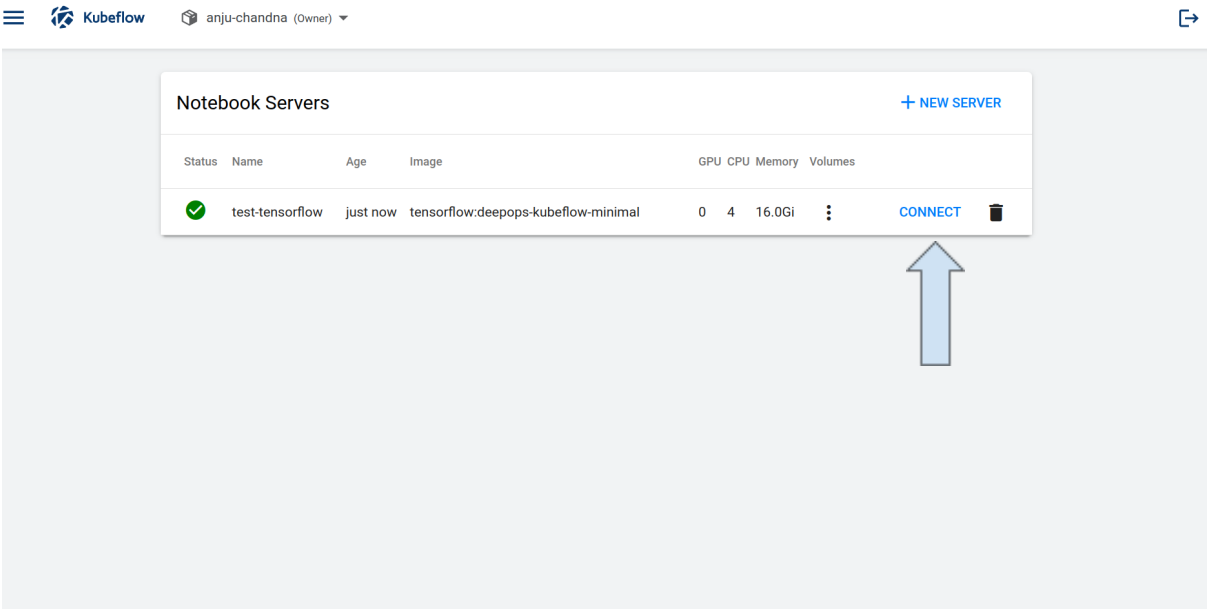

Note - When we select NVIDIA MIG GPU (20GB or less) then we will see GPU count 0 in the Notebook Servers. This is a bug in Kubeflow and can be ignored.

DO NOT FORGET TO DELETE YOUR NOTEBOOK SERVERS AFTER COMPLETING YOUR TASK TO AVOID UNNECESSARY BLOCKING OF RESOURCES!